

Azure Trailblazer Academy

Azure Data Factory (ADF) Lab

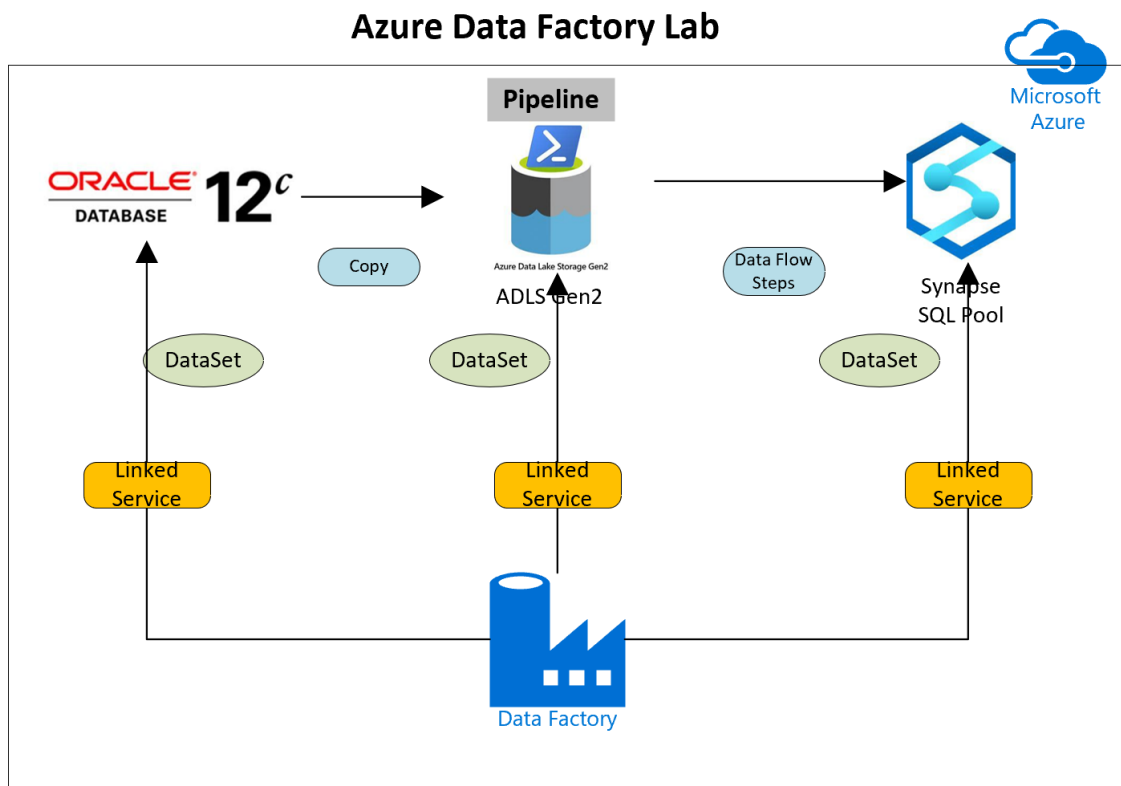
Overview

Azure Data Factory is a PaaS cloud-based ETL & data integration tool that allows you to create data-driven workflows in the cloud for orchestrating and automating data movement and data transformation.

It is a great tool to orchestrate big data and operationalize processes to refine these enormous stores of raw data into actionable business insights.

Lab Overview

This lab will help you gain the experience to ingest data from on-premises databases such as Oracle, SAP, Teradata, Hortonworks, DB2, SQL Server and Cloudera to Azure Data storage, databases and data warehouses services. It will showcase the steps to build a pipeline using ADF to ingest the data into ADLS GEN2 storage and secure the PII data using data transformation functions using the Data Flow activity and finally store the data in Synapse SQL Pool (Data warehouse) for building BI dashboards.



Pre-requisites

- Write Access to Azure Data Lake Storage Account (ADLS Gen2)
- Read Access to Sample HR schema in Oracle Database
- Write Access to Synapse SQL Pool Data warehouse

Automated Deployment

This automated deployment script will create the required Azure services for this lab. Press the "Deploy to Azure" button below.



https://portal.azure.com/#create/Microsoft.Template/uri/https%3A%2F%2Fraw.githubusercontent.com%2Fmicrosoft%2FAzureTrailblazerAcademy%2Fmaster%2Fmonth2%2Fflabs%2Fflab_adf%2Fscripts%2Fflab2_data_deployment.json

Enter the following information

- Subscription: Enter your subscription.
- Resource group: Select 'Create new' under Resource group. Enter 'ata-adf-lab-<YourName>-rg'
- Region: Select 'East US'.
- SQL_Server Name: Enter 'ata-adf-lab-sql-<YourName>'
- Server_location: Enter 'eastus'

Custom deployment

Deploy from a custom template

Subscription * ⓘ

Resource group * ⓘ [Create new](#)

Parameters

Region * ⓘ

SQL_Server Name ⓘ

Server_Location ⓘ

Administrator Login * ⓘ ✓

Administrator Login Password ⓘ ✓

Collation ⓘ

Database Name * ⓘ ✓

Sku Name ⓘ




Storage Accounts_Data Lake Gen2 ⓘ ✓

Datalakestorage_location ⓘ

[Review + create](#) [< Previous](#) [Next : Review + create >](#)

- AdministratorLogin: 'azureadmin'
- Administrator Login Password: 'Ataadf123!'
- Database Name: 'ataadflabsqldb'
- Storage_Accounts_Data Lake Gen2: 'ataadflabstorage<YourName>/'
- Click on 'Review and Create'
- Click on 'Create'

Please check if it created the following services after successful deployment.

<input type="checkbox"/>	 ataadflabsqldb (ata-adf-lab-sql-srini/ataadflabsqldb)	Synapse SQL pool (data warehouse)
<input type="checkbox"/>	 ata-adf-lab-sql-srini	SQL server
<input type="checkbox"/>	 ataadflabstoragesrini	Storage account

Temporary access to an Oracle database for completing this lab will be provided in the class.

Task List

- [Task-1: Create Azure Data Factory Service](#)
- [Task-2: Create linked services](#)
- [Task-3: Create Copy Activity to ingest Oracle HR Employee data to Azure Storage](#)
- [Task-4: Create Data Flows to Transform Employee data and store it in Synapse SQL Analytics](#)

- [Task-5: Build a pipeline to connect Copy and Dataflow activities and Test the Execution](#)

Task-1: Create Azure Data Factory Service

1. Provision an ADF service
 - Type 'Data factories' in the search bar.
 - Select 'Data factories' and select 'add' to create a new service

Provide the following info:

- Name: 'ata-adf-lab-<YourName>'
- Subscription: Make sure to select your subscription
- Select 'ata-adf-lab-<YourName>', the resource group you have created with custom deployment.
- Location: select 'East US'
- Enable GIT: uncheck the box

Click on 'Create' button

New data factory

Name *

ata-adf-lab-srin

Version ⓘ

V2

Subscription *

Resource Group *

ata-adf-lab-srini-rg

[Create new](#)

Location * ⓘ

East US

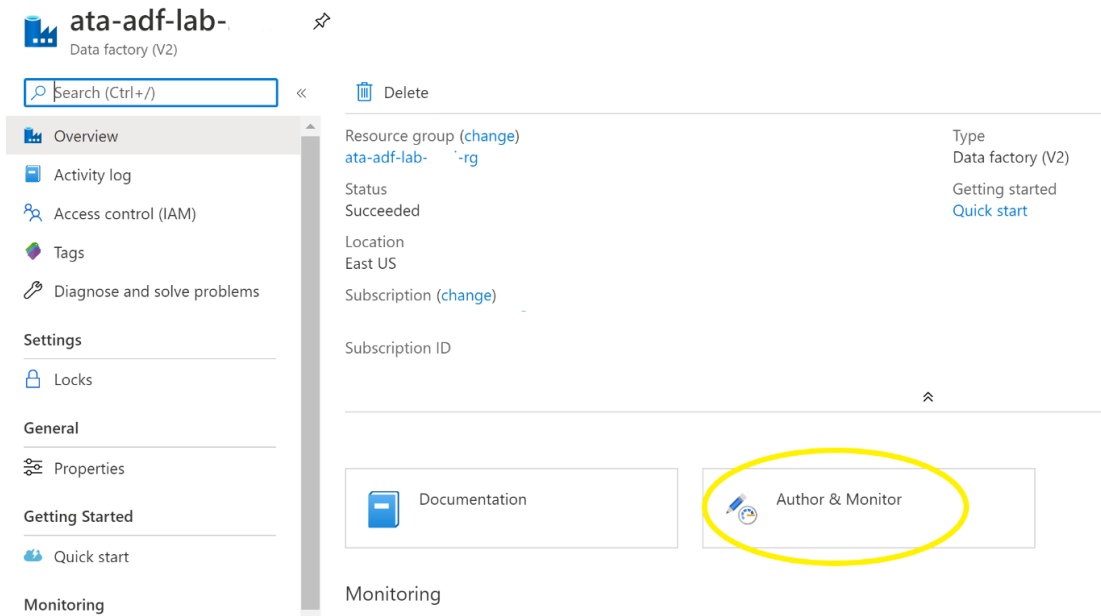
Enable GIT ⓘ

☐


Create

2. Open ADF Author Tool

- Select 'Go to resource' when it completes the deployment.
- Select 'Author & Monitor' in the middle of the screen.




Opens up a new tab introducing the drag and drop interface to build pipelines.



Turn on Dataflow Debug functionality to test the modifications.

- To preview data, we need to turn on the dataflow debug option

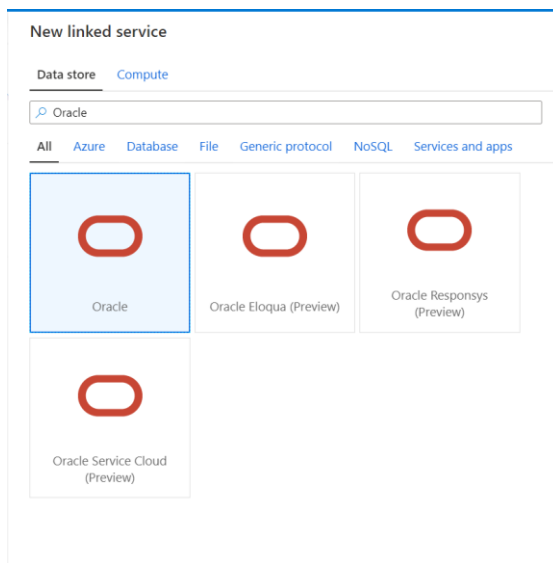


Task-2: Create linked services

You will be creating connection linked services to sync and source systems such as Oracle, ADLS Gen2 and Synapse in this task.

1. Create Oracle Linked Service

- Select Management hub (ToolBox) icon on the left and select 'Linked Services' under Connections section.
- Select 'New' under Linked Services
- Search for 'Oracle' under 'Data Store' and select 'Oracle Database' and click on 'Continue'



Enter the following Oracle Connect Info:

- Name: Enter 'OracleDB12cHR'
- Leave the default 'Connection string' option
- Host: Enter Instructor provided server IP address
- Port: 1521 (Default Oracle Port)
- Connection type: Select 'Oracle Service Name'
- Service name: **'nonpdb'**
- Username: hr
- Password: hr

New linked service (Oracle)

Name *

OracleDB12cHR

Description

Connect via integration runtime *

AutoResolveIntegrationRuntime



Connection string

Azure Key Vault

Host *

Port

1521

Connection type

Oracle Service Name



Service name *

pdb1

User name *

hr

Password

Azure Key Vault

Password *

..

Create

Back



Connection successful



Test connection

Cancel

- Select 'test connection' to verify the successful connection

Click on 'Create' button to create the Oracle linked service.

2. Create ADLS Gen2 Storage Linked Service

- Select 'New' under Linked Services
- Search for 'Gen2' under data store and select ADLS Gen2 Click on 'Continue' and enter the following information:

- Name: Enter 'Gen2HRStorage'
- Authentication Method: Leave the default 'Account Key' selection
- Account Selection method: Leave the default 'From Azure subscription'
- Storage account name: select 'ataadflabstorage<YourName>'
- Click on 'Test connection' Click on 'Create' after successful connection to create the ADLS Gen2 storage linked service

New linked service (Azure Data Lake Storage Gen2)

Hosted integration runtime is higher than version 4.0 if connecting via self-hosted integration runtime.

Name *

Gen2HRStorage

Description

Connect via integration runtime *

AutoResolveIntegrationRuntime

Authentication method

Account key

Account selection method

☒ From Azure subscription ☐ Enter manually

Azure subscription

Storage account name *

ataadflabstoragesrini

Test connection

☒ To linked service ☐ To file path

Annotations

+ New

Create

Back



Connection successful



Test connection

Cancel

3. Create Synapse Analytics Linked Service

- Select 'New' under Linked Services
- Search for 'Synapse' under data store and select 'Azure Synapse Analytics'
- Click on 'Continue'

Enter the following info:

- Name: Enter 'SynapseDBHR'
- Server name: select 'ata-adf-lab-sql-<yourname>'
- Database name: select 'ataadflabsqldb'
- Username: enter 'azureadmin'
- Password: enter the default 'ataadf123!'
- Click on 'test connect' to test the connection.

Click on 'Create' after the successful connection to create the Synapse linked service

New linked service (Azure Synapse Analytics (formerly SQL DW))

Name *

SynapseDBHR

Description

Connect via integration runtime *

AutoResolveIntegrationRuntime

Connection string

Azure Key Vault

Account selection method

☒ From Azure subscription ☐ Enter manually

Azure subscription

Server name

ata-adf-lab-sql-srini

Database name

ataadflabsqldb

Authentication type *

SQL authentication

User name *

azureadmin


Password

Azure Key Vault

Password *

Create

Back



Test connection

Cancel


You have successfully created connection linked services to Oracle, ADLS Gen2 and Synapse SQL Pool.

Task-3: Create Copy Activity to ingest Oracle HR Employee data to Azure Storage

We have established the connection services to the source Oracle DB and the sink Azure storage, you will create a copy activity to ingest the data from Oracle to Azure Storage in this task.

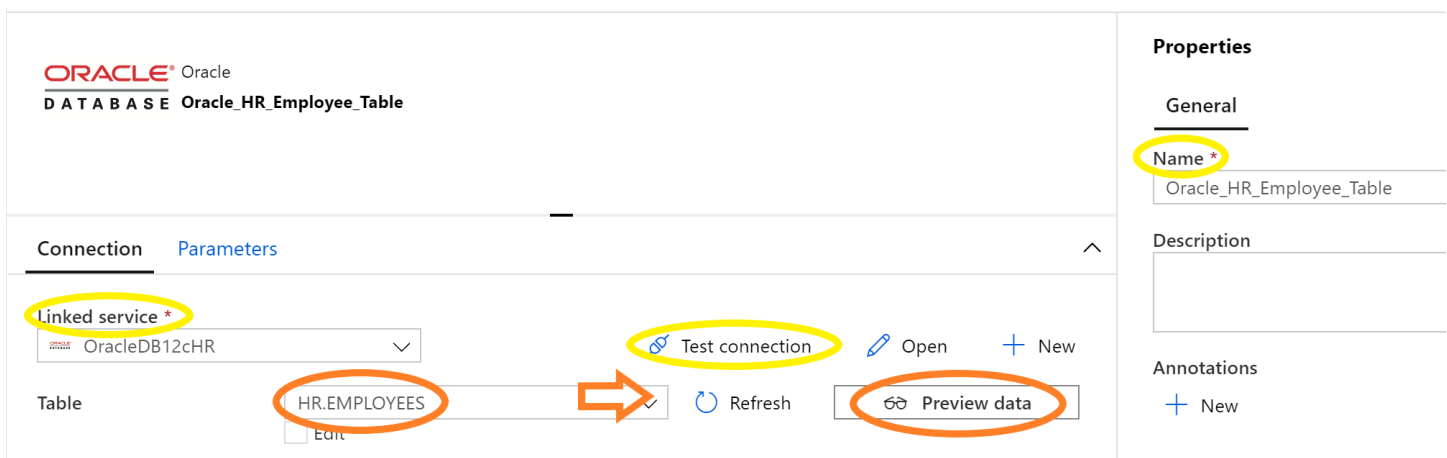
1. Select 'Pencil' icon on the left and select three dots next to pipelines to select 'new pipeline' action.

- Name the pipeline as 'IngestOracleHREmployeeData' under the properties section on the right side.
- Drag the 'Copy data' from 'Move & transform' section under 'Activities' list to the canvas in the middle of the screen.
- Name the copy activity as 'CopyOracleHREmpToGen2' below the canvas under the 'General' tab



2. Select 'Source' tab next to 'General' to define the source system.

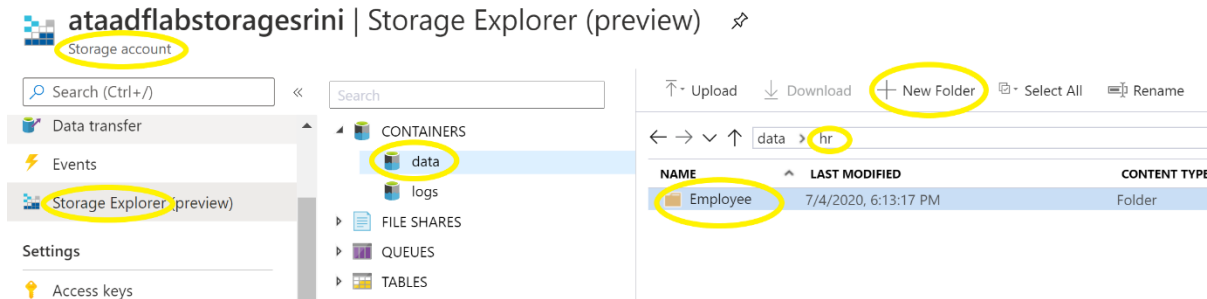
- Select 'New' to create a new source dataset
- Select 'Oracle database' as the data store after filtering with 'Oracle' and click on 'Continue'
- Select 'Open' to define the source dataset
- Name the dataset as 'Oracle_HR_Employee_table' under the properties section on the right side.
- Select the 'OracleDB12cHR' linked service and click on 'Test connection' to test the connectivity.
- Filter table list by typing 'hr.emp' and select 'HR.EMPLOYEES' table and select 'Preview data'.



Make sure you are able to see the employee folder

You will have to enter the file system and directory path

- First, we need to create the directory path in the storage account. Leave this browser tab as is and switch to the Azure Services tab in the browser.
- Access storage account and open up the 'storage explorer' to create 'hr' under 'data' file system and create 'employee' as a subfolder under 'hr' folder.**



3. Select 'Sink' tab next to 'Source' to define the sink system

- Select 'New' to create a new sink data set
- Select 'ADLS Gen2' as the data store after filtering with 'Gen2' and click on 'Continue'
- Select 'Open' to define the sink dataset
- Select 'Delimited Text' as the format and click on 'Continue'
- Name the dataset as 'Gen2HREmpData'
- Select the 'Gen2HRStorage' linked service.

Set properties

Name

Gen2HREmpData

Linked service *

Gen2HRStorage

File path

data

/ hr/employee

/ File

First row as header

☐

Import schema

☒ From connection/store ☐ From sample file ☐ None

▶ Advanced

OK

Back

Cancel

- Select 'Gen2HREmpData' under Datasets section on the left to set the first row as the header.
- Check the box

The screenshot shows the 'Set properties' dialog for a dataset named 'Gen2HREmpData'. The 'Name' field is 'Gen2HREmpData'. The 'Linked service' is 'Gen2HRStorage'. The 'File path' is 'data / hr/employee / File'. The 'First row as header' checkbox is checked. The 'Import schema' section shows 'From connection/store' selected. The 'Advanced' section is collapsed. The 'OK' button is highlighted.

5. Test the copy activity by clicking on 'Debug' option just above the canvas.

- It will start the process and put it in the queue. Wait till it finishes.
- Check for the status and make sure it is successful

Copy data

OracleHREmpToGen2

Parameters Variables **Output**

Pipeline run ID: 798c8f61-d66a-4a7a-966a-82a43d159576 [@] View debug run consumption

NAME	TYPE	RUN START	DURATION	STATUS	INTEGRATION RUNT
OracleHREmpToGen2	Copy	2020-07-04T22:44:14.556	00:00:06	✓ Succeeded	DefaultIntegratic

6. Verify the data ingestion in ADLS Gen2 storage.

Switch to Azure services tab and access the storage account.

- Open up the 'Storage Explorer' and access the data file system and drill down to 'hr' and 'Employee' folder.
- Confirm the 'HR.EMPLOYEES.txt' file. Double click on the file to download and view the data.

You can see how the phone numbers are ingested as text. This is PII data and we should protect this data.

ataadflabstoragesrini | Storage Explorer (preview)

Storage account

Search (Ctrl+/)

CONTAINERS

- data
- logs

FILE SHARES

HR.EMPLOYEES (1).txt - Notepad

File Edit Format View Help

EMPLOYEE_ID,FIRST_NAME, LAST_NAME, EMAIL, PHONE_NUMBER, HIRE_DATE, JOB_ID, SALARY, COMMISSION_PCT, MANAGER_ID, DEPARTMENT_ID

100, "Steven", "King", "SKING", 515.123.4567, 2003-06-17 00:00:00.000000, "AD_PRES", 24000.00, , 90

101, "Neena", "Kochhar", "NKOCHHAR", 515.123.4568, 2005-09-21 00:00:00.000000, "AD_VP", 17000.00, , 100, 90

102, "Lex", "De Haan", "LDEHAAN", 515.123.4569, 2001-01-13 00:00:00.000000, "AD_VP", 17000.00, , 100, 90

103, "Alexander", "Hunold", "AHUNOLD", 590.423.4567, 2006-01-03 00:00:00.000000, "IT_PROG", 9000.00, , 102, 60

104, "Bruce", "Ernst", "BERNST", 590.423.4568, 2007-05-21 00:00:00.000000, "IT_PROG", 6000.00, , 103, 60

Task-4: Create Data flow to transform Employee data and store it in Synapse SQL Analytics

We have noticed the PII data we just ingested into ADLS Gen2 storage. You will secure the PII data using the Data flow functionality in this task.

1. Create Source data

- Drag and drop the data flow activity into the canvas from the 'Move & Transform' section.
- Select 'create new data flow' and select 'Data Flow' option and click on 'OK'.

Adding data flow

☐ Use existing data flow ☒ Create new data flow



Mapping Data Flow

Code free data transformation at scale



Wrangling Data Flow (Preview)

Code free data preparation at scale

OK

Cancel

- Select "Source" data
- Name the data flow as 'SecurePIIData'.

Provide the following information under 'Source settings' tab:

- Output stream name: Gen2HrEmpDataOut
- Source type - Select 'Dataset'
- Dataset - Select 'Gen2HREmpData'

✓ Validate

Gen2HREmpDataOut
Columns: 0 total

Source settings Source options Projection Optimize Inspect Data preview Description

Output stream name * Gen2HREmpDataOut [Learn more](#)

Source type * Dataset

Dataset * Gen2HREmpData

Test connection Open + New

Options

- ☒ Allow schema drift ⓘ
- ☐ Infer drifted column types ⓘ
- ☐ Validate schema ⓘ

Skip line count

Sampling * ☐ Enable ☒ Disable ⓘ

Properties

General

Name * SecurePIIData

Description

- To preview data, we need to turn on the dataflow debug option.
- Click on 'Import projections' in the 'Projection' tab to view all the columns.

✓ Validate all Refresh Discard all Data flow debug ARM template

✓ Validate Debug Settings

Gen2HREmpDataOut
Columns: 11 total

Source settings Source options Projection Optimize Inspect Data preview Description

Define default format Detect data type Import projection Reset schema

Column name	Type	Format
EMPLOYEE_ID	12s short	Specify format
FIRST_NAME	abc string	Specify format
LAST_NAME	abc string	Specify format
EMAIL	abc string	Specify format
PHONE_NUMBER	abc string	Specify format
HIRE_DATE	abc string	Specify format
JOB_ID	abc string	Specify format
SALARY	1.2 double	Specify format
COMMISSION_PCT	1.2 double	Specify format
MANAGER_ID	12s short	Specify format
DEPARTMENT_ID	12s short	Specify format

2. Create a data flow step to hash the first 6 digits of the phone number.

- Add another data flow step to secure the phone number by selecting '+' sign and select 'Derived Column' under 'Schema modifier' section.

The screenshot shows the 'Gen2HREmpData' data flow step. A search bar is active, and the 'Derived Column' option is highlighted under the 'Schema modifier' section. The data preview table is visible below.

LAST_NAME	EMAIL	PHONE_NUMBER
King	SKING	515.123.4567
Kochhar	NKOCHHAR	515.123.4568
De Haan	LDEHAAN	515.123.4569
Hunold	AHUNOLD	590.423.4567
Ernst	BERNST	590.423.4568
Austin	DAUSTIN	590.423.4569
Valli	VPATABAL	590.423.4560

- Name the data flow as 'HREmpSecurePhoneNumber'
- Select 'PHONE_NUMBER' column from the list
- Click on 'Enter Expression' box

The screenshot shows the 'Derived column's settings' panel. The 'Output stream name' is 'HREmpSecurePhoneNumber', the 'Incoming stream' is 'Gen2HREmpData', and the 'Columns' list contains 'PHONE_NUMBER'. The 'Enter expression...' box is highlighted with a yellow arrow.

Write RegExpression to replace the digits before the '.' with '#' as shown in the picture:

```
regexReplace({PHONE_NUMBER}, '\\d+\\. ', '###.')
```

- Click on 'Refresh' to verify the output of the function.
- Click 'Save and finish' button after you satisfy the result

Visual expression builder

Expression reference documentation

OUTPUT SCHEMA << FUNCTIONS << EXPRESSION FOR FIELD "PHONE_NUMBER" Save

ANY PHONE_NUMBER

Filter...

regexReplace({PHONE_NUMBER}, '\\d+\\. ', '###.')

123 EMPLOYEE_ID
abc FIRST_NAME
abc LAST_NAME
abc EMAIL
abc PHONE_NUMBER
abc HIRE_DATE

Data preview

Output: PHONE_NUMBER abc

PHONE_NUMBER abc
###.###.4567
###.###.4568
###.###.4569

Refresh

Save and finish Cancel Clear contents

3. Create a data flow step to store the transformed data into Synapse SQL Pool.

We are now ready to output the transformed data into Synapse SQL Pool

- Select '+' and select 'Sink' in the 'Destination' section. Last one in the list.
- Enter 'sinkToSynapseDb' for Output stream name
- Click the '+' next to Dataset to create new Azure Synapse Analytics (formerly SQL DW) dataset

IngestOracleHREmp... SecurePIIData

Validate Debug Settings

Gen2HrEmpDataOut

Import data from Gen2HREmpData

Reference: 1
Columns: 11 total

sinkToSynapseDb

Columns: 11 total

Add Source

Sink Settings Mapping Optimize Inspect Data preview

Output stream name * sinkToSynapseDb Learn more

Incoming stream * HREmpSecurePhoneNumber

Sink type * Dataset

Dataset * Select... + New

Options

☒ Allow schema drift

☐ Validate schema

New dataset

In pipeline activities and data flows, reference a dataset to specify the location and structure of your data within a data store. [Learn more](#)

Select a data store

All

Azure


Database


File


Generic protocol


NoSQL


Services and apps



Azure Blob Storage



Azure Cosmos DB (SQL API)



Azure Data Lake Storage Gen1



Azure Data Lake Storage Gen2


Azure SQL Database


Azure SQL Database Managed Instance


Azure Synapse Analytics (formerly SQL DW)


Snowflake


Amazon Marketplace Web Service

- For dataset name, enter 'Synapse_HR_Employee_table'
- Check the 'Edit' and enter schema: dbo and table name: employee
- Click 'OK'

Set properties

Name

Linked service *

☒ Select from existing table ☐ Create new table

Table name

☒ Edit

Import schema

☒ From connection/store ☐ None

▶ Advanced

- Click the 'Settings' tab
- Check the 'Recreate table' radio button
- Uncheck the 'Enable staging' checkbox

IngestOracleHREmp...SecurePIIDataSynapse_HR_Employe...

✓ Validate ⚙ Debug Settings

Gen2HrEmpDataOut

Import data from Gen2HREmpData

+

HREmpSecurePhoneNu...

Creating/updating the columns
EMPLOYEE_ID, FIRST_NAME,
LAST_NAME, EMAIL,
PHONE_NUMBER, HIRE_DATE,

+

sinkToSynapseDb

Columns:
11 total

Add Source

SinkSettingsMappingOptimizeInspectData preview ●

📘 We recommend enabling staging to improve performance with Azure Synapse Analytics datasets.

Update method

☒ Allow insert
☐ Allow delete
☐ Allow upsert
☐ Allow update

Table action

☐ None ☒ Recreate table ☐ Truncate table

Enable staging

☐

Batch size

ⓘ

Pre SQL scripts

Post SQL scripts

Task-5: Build a pipeline to connect Copy and Data flow activities and Test the Execution

We have a create copy activity to ingest the data from Oracle to Azure Storage and Dataflow activity to secure PII data. You will connect both activities to build a pipeline to execute the end to end functionality.

1. Connect Copy activity with Data Flow activity
 2. Trigger the pipeline execution.
- Click on 'Debug' to test the data transformation to secure the PII data and storing it to Synapse SQL Analytics database.

The screenshot shows the Azure Data Factory 'Activities' pane on the left, with 'Copy data' and 'Data flow' selected. The main canvas displays a pipeline diagram with a 'Copy data' activity (labeled 'CopyOracleHREmpToGen2') connected to a 'Data flow' activity (labeled 'SecurePIIData'). The 'Data flow' activity is highlighted with a red circle. The top bar shows the pipeline name 'IngestOracleHREmp...' and the 'Debug' button is visible.

- Execution will be queued and wait till it finishes.

The screenshot shows the 'Output' tab of the pipeline execution. The pipeline run ID is 645cc886-28e6-4fc2-952b-d1b0518492ff. The output table shows the following data:

NAME	TYPE	RUN START	DURATION	STATUS	INTEGRATION RUNTIME
SecurePIIData	ExecuteDataFlk	2020-07-05T23:16:28.182	00:01:03	Succeeded	DefaultIntegrationRunti
OracleHREmpToGen2	Copy	2020-07-05T23:16:20.674	00:00:06	Succeeded	DefaultIntegrationRunti

The 'SecurePIIData' and 'OracleHREmpToGen2' rows are circled in yellow. The 'Status' column shows 'Succeeded' for both activities, which are also circled in red.

4. Verify the transformed data was inserted into Azure Synapse Analytics SQL Pool.

- Select 'Firewalls and virtual networks' from 'Security' section
- Click the '+ Add client IP'
- Click on 'Save' button after you added the rule.

Home > Resource groups > ata-adf-lab-bpham > ata-adf-lab-sql-bpham

ata-adf-lab-sql-bpham | Firewalls and virtual networks

SQL server

Search (Ctrl+/)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Settings

Quick start

Failover groups

Manage Backups

Active Directory admin

SQL databases

SQL elastic pools

Deleted databases

Import/Export history

DTU quota

Properties

Locks

Security

Auditing

Firewalls and virtual networks

Private endpoint connections

Security Center

Transparent data encryption

Intelligent Performance

Save Discard + Add client IP

Deny public network access Yes No

To set Deny Public Network Access, click here to create a new private endpoint.

Minimum TLS Version > 1.0 > 1.1 > 1.2

You are setting the Minimal TLS Version property for all SQL Database and SQL Data Warehouse databases associated with the server. Any login attempts from clients using TLS version less than the Minimal TLS Version shall be rejected.

Connection Policy Default Proxy Redirect

Allow Azure services and resources to access this server Yes No

Connections from the IPs specified below provides access to all the databases in ata-adf-lab-sql-bpham.

Client IP address 68.170.94.159


Rule name	Start IP	End IP
ClientIPAddress_2020-10...	68.170.94.159	68.170.94.159

Virtual networks + Add existing virtual network + Create new virtual network

Rule name	Virtual network	Subnet	Address Range	Endpoint
No vnet rules for this server.				

- Switch to Azure Services and access the Synapse SQL Analytics to verify the employee data with secured phone numbers
- Select the Overview and scroll down to select the database
- Select the 'Query editor' and enter the login info again
- Login with 'azureadmin' and password 'Ataadf123!' if needed.
- With successful login, you will see the tables and views.
- Access synapse SQL Analytics 'ataadflabsqldb' and select 'Query editor' from the 'Common Tasks' section.

Home > Resource groups > ata-adf-lab-bpham > ata-adf-lab-sql-bpham > ataadflabsqldb (ata-adf-lab-sql-bpham/ataadflabsqldb)

ataadflabsqldb (ata-adf-lab-sql-bpham/ataadflabsqldb) | Query editor (preview) 

Synapse SQL pool (data warehouse)

Search (Ctrl+/) << Login + New Query ↑ Open query Feedback

Overview

Activity log

Tags

Diagnose and solve problems

Settings

Workload management

Maintenance schedule

Quick start

Geo-backup policy

Connection strings

Properties

Locks

Security

Auditing

Data Discovery & Classification

Dynamic Data Masking

Firewalls and virtual networks

Security Center

Transparent data encryption

Common Tasks

View streaming jobs


Load Data

Query editor (preview) ✓

Build dashboards + reports

Model and cache data

Open in Visual Studio



Welcome to SQL Database Query Editor

SQL server authentication

Login *

azureadmin

Password *

OK

Active Directory authentication

Continue as brpham@microsoft.com

OR

✓

- Access synapse SQL Analytics 'ataadflabsqldb' and select 'Query editor' from the 'Common Tasks' section.
- Enter 'select * from dbo.employee' to view the ingested data from the Oracle database.
- You should see all the employee data with the secured phone number.

ataadflabsqldb (ata-adf-lab-sql-srini/ataadflabsqldb) | Query editor (preview)

Synapse SQL pool (data warehouse)

Search (Ctrl+/) Login + New Query ↑ Open query Feedback

Security

- Advanced data security
- Auditing
- Firewalls and virtual networks
- Transparent data encryption

Common Tasks

- View streaming jobs
- Load Data
- Query editor (preview)**
- Build dashboards + reports
- Model and cache data
- Open in Visual Studio

Monitoring

- Query activity
- Alerts
- Metrics

ataadflabsqldb (azureadmin)

Showing limited object explorer here. For full capability please open SSDT.

- Tables
- Views
- Stored Procedures

Query 1

Run Cancel query Save query Export data as Show only Editor

1 select * from [HR].[Employee]

Results Messages

Search to filter items...

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE
100	Steven	King	SKING	###.###.4567	2003-06-17 00:00:0
101	Neena	Kochhar	NKOCHHAR	###.###.4568	2005-09-21 00:00:0
102	Lex	De Haan	LDEHAAN	###.###.4569	2001-01-13 00:00:0
103	Alexander	Hunold	AHUNOLD	###.###.4567	2006-01-03 00:00:0

Congratulations!! You have successfully ingested the data from Oracle to Azure Datalake, transformed the data as per your needs such as securing the data and stored the data into Azure Synapse Analytics for your business analytics needs!!!